

Part II. Process Management

- Chapter 3 Processes
- Chapter 4 Threads
- Chapter 5 CPU Scheduling
- Chapter 6 Process Synchronization
- Chapter 7 Deadlocks

Chapter 3. Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Inter-Process Communication (IPC)

1. Process Concept

- An operating system executes a variety of programs:
 - Batch system – *jobs*
 - Time-shared systems – user programs or *tasks*
- Textbook uses the terms *job* and *process* almost interchangeably.
- **Process**
 - An instance of execution of a program that is loadable into the main memory.
 - Process execution must progress in sequential fashion.
- A process includes:
 - **Text section**
 - **CPU status** – **Program counter** and **registers**
 - **Stack** – for local variables
 - **Data section** – for global variables
 - **Heap** – dynamically allocated memory
 - See Figure 3.1.

Process State

- As a process executes, it changes *state*.
 - **new**: The process is being created.
 - **ready**: The process is waiting to be assigned to a processor.
 - **running**: Instructions are being executed.
 - **waiting**: The process is waiting for some event to occur, e.g., i/o completion and reception of a signal.
 - **terminated**: The process has finished execution.
 - See Figure 3.2.
 - ☺ *Where is a process when it is in the ready state, or in the running state?*

Process Control Block (PCB)

- The operating system must know specific information about processes in order to manage and control them.
- **Process state information**
 - CPU registers
 - Program counter
- **Process control information**
 - CPU scheduling information
 - I/O status information
 - Memory-management information
 - Accounting information

Process Control Block - continued

- This collection of process information is kept in and accessed through the **Process control block (PCB)**, also called the task control block (TCB).
 - Process state – new, ready, ...
 - Program counter
 - CPU registers
 - CPU scheduling information – priority, pointers to queues, ...
 - Memory-management information – base and limit registers, addresses of ...
 - Accounting information – the amount of CPU, time limits, account numbers, ...
 - I/O status information – allocated i/o devices, open files, ...
 - See Figure 3.3.
 - ☺ *Where are they located?*
- ☺ *What is a PCB?*

CPU Switch From Process to Process

- See Figure 3.4.
 - Operating system, especially process scheduler, that could be invoked from an interrupt service routine.
 - Interrupt from i/o devices or system call from any other process

2. Process Scheduling

- ☺ *Where are PCBs?*
 - Queues in main memory
- ***Job queue*** – set of all processes in the system
- ***Ready queue*** – set of all processes residing in main memory, ready and waiting to execute
- ***Device queues*** – set of processes waiting for an I/O device
- Process migration between the various queues, by **process scheduler**
- See Figure 3.6.
- See Figure 3.7.

Schedulers

- *Long-term scheduler* (or **job scheduler**) – selects which processes should be brought into the ready queue
- *Short-term scheduler* (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
- See Figure 3.8.

Schedulers - continued

- Short-term scheduler is invoked **very frequently** (milliseconds) \Rightarrow (must be fast).
- Long-term scheduler is invoked very **infrequently** (seconds, minutes) \Rightarrow (may be slow).
- The long-term scheduler controls the *degree of multiprogramming* (the number of processes in main memory).
 - Processes can be described as either:
 - *I/O-bound process* – spends more time doing I/O than computations, many short CPU bursts
 - *CPU-bound process* – spends more time doing computations; few very long CPU bursts
 - ☺ *Which one first?*
 - ☺ *How to mix them?*

Context Switch

- Context of a process – PCB
- **Context switch** – When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process.
- Context-switch time is overhead; the system does no useful work while switching.
- Time dependent on hardware support
 - E.g., several set of registers => switching of register sets => very fast

3. Operations on Processes

- Process creation
- Process termination

Process Creation

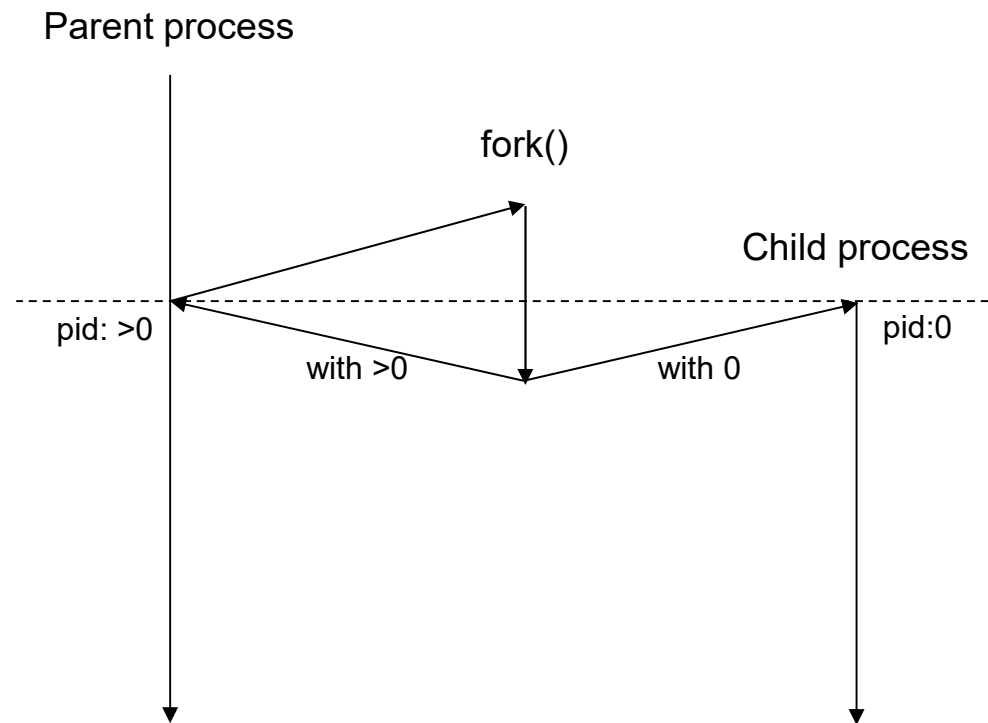
- **Parent process** creates **children processes**, which, in turn create other processes, forming a tree of processes.
- ☺ *What does creation of a process mean ???*
 - Creation of a new PCB, and put it into the job queue
- ☺ *How can a parent process create a child process???*
 - System call
- Resource sharing
 - Parent and children share all resources, or
 - Children share subset of parent's resources, or
 - Parent and child share no resources
- Execution
 - Parent and children execute concurrently, or
 - Parent waits until children terminate
- Address space
 - Child duplicate of parent, or
 - Child has a program loaded into it

Process Creation - continued

- ☺ *How to create?*
- UNIX examples
 - **fork()** system call creates a new child process
 - Same text section
 - Different data section, stack, heap
 - Different CPU status
 - **exec()** system call used after a **fork** to replace the process' memory space with a new program
 - Now, different text section
 - Different data section, stack, heap
 - Different CPU status
 - **wait()** system call used to move itself off the ready queue until the termination of the child
 - See Figure 3.10 – 3.11.

C Program Forking Separate Process

- `pid = fork();`



Process Termination

- ☺ *How to terminate?*
- Process executes last statement and asks the operating system to decide it (**exit**).
 - Output data from child to parent (via **wait**)
 - Process' resources including PCB are deallocated by operating system.
- Parent may terminate execution of children processes (**abort**).
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - If parent is exiting
 - Some operating system do not allow child to continue if its parent terminates.
 - All children terminated - *cascading termination*

4. Cooperating Processes

Cooperating Processes

- *Independent* process cannot affect or be affected by the execution of another process
- *Cooperating* process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

In the game, Tetris,
One process for dropping blocks
Another process for turning blocks
☺ *How can they operate?*

Producer-Consumer Problem

- Paradigm for cooperating processes
 - *Producer* process produces information that is consumed by a *consumer* process.
 - Need a buffer that can be filled by the producer and emptied by the consumer.
 - *Unbounded-buffer* places no practical limit on the size of the buffer.
 - *Bounded-buffer* assumes that there is a fixed buffer size.
 - What kind of data type would be proper?
 - The producer and consumer must be synchronized.

5. Interprocess Communication (IPC)

Interprocess Communication (IPC)


- Mechanism for processes to communicate and to synchronize their actions
- Message-passing system – processes communicate with each other without resorting to shared variables
 - IPC facility provides two operations:
 - **send**(*message*) – message size fixed or variable
 - **receive**(*message*)
 - If *P* and *Q* wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via `send()`/`receive()`
 - Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)

Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?
- Two communication models
 - See Figure 3.14.


Direct Communication

- Concept of **message passing**
- Processes must name each other explicitly:
 - **send** (P , *message*) – send a message to process P
 - **receive** (Q , *message*) – receive a message from process Q
- Properties of communication link
 - Links are established automatically.
 - A link is associated with exactly one pair of communicating processes.
 - Between each pair there exists exactly one link.
 - The link may be unidirectional, but is usually bi-directional.



Indirect Communication

- Concept of **shared memory**
- Messages are directed and received from **mailboxes** (also referred to as **ports**).
 - Each mailbox has a unique id.
 - Processes can communicate only if they share a mailbox.
- Properties of communication link
 - Link established only if processes share a common mailbox.
 - A link may be associated with many processes.
 - Each pair of processes may share several communication links.
 - Link may be unidirectional or bi-directional.



Processes

Indirect Communication - continued

- Operations:
 - create a new mailbox.
 - send and receive messages through mailbox.
 - destroy a mailbox.
- Primitives are defined as:
 - send**(A , $message$) – send a message to mailbox A
 - receive**(A , $message$) – receive a message from mailbox A

Indirect Communication - continued

- Mailbox sharing:
 - P_1 , P_2 , and P_3 share mailbox A.
 - P_1 sends; P_2 and P_3 receive.
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes.
 - Allow only one process at a time to execute a receive operation.
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Synchronization

- Message passing may be either **blocking** or **non-blocking**.
- **Blocking** is considered *synchronous*.
 - **Blocking send** has the sender block until the message is received.
 - **Blocking receive** has the receiver block until a message is available.
 - Also called *Rendezvous*
- **Non-blocking** is considered *asynchronous*.
 - **Non-blocking send** has the sender send the message and continue.
 - **Non-blocking receive** has the receiver receive a valid message or null.

In the game, Tetris,
One process for dropping blocks
Another process for turning blocks
☺ *Synchronous or asynchronous?*

Buffering

- Queue of messages attached to the link; implemented in one of three ways:
 1. Zero capacity – 0 messages
Sender must wait for receiver (rendezvous).
 2. Bounded capacity – finite length of n messages
Sender must wait if link full.
 3. Unbounded capacity – infinite length
Sender never waits.

Other IPC Mechanisms

- UNIX

- Pipes – FIFO

- `$ cat /etc/passwd | grep mlee`

- Signals

- To notify a process of an event
 - Asynchronous
 - By one process to another, including itself; by the kernel to a process
 - `kill(int pid, int sig)`
 - `raise(int sig)`
 - `signal(int sig, SIGARG func)`
 - Can ignore or catch the signal, depending on func

