

Chapter 11. File System Implementation

- File System Structure
- File System Implementation
- Directory Implementation
- Allocation Methods
- Free-Space Management
- Efficiency and Performance
- Recovery

1. File-System Structure

- File structure
 - Logical storage unit
 - Collection of related information
- File system resides on secondary storage (disks)
- File system organized into layers.

Layered File System

- See Figure 11.1.

2. File-System Implementation

- On disk, each file system
 - **Boot control block**, or called **boot block** in UFS, or called **partition boot sector** in NTFS
 - First block of a partition
 - **Partition control block**
 - **Superblock** in UFS
 - **Master file block** in NTFS
 - Directory structure
 - To organize the files
 - **File control block**
 - Could be stored in directory structure, or
 - Independently
 - **Inode** in UFS
 - Stored within the MFB in NTFS, which uses a relational database structure
 - See Figure 11.2.

- In memory, to improve performance:
 - Partition control block for each mounted partition
 - Directory structure for recently accessed directories
 - File descriptor tables
 - Open file table
 - Inode table

- See Figure 11.3.

Virtual File Systems

- How to support concurrently multiple types of file system
- **Virtual File Systems (VFS)** provide an object-oriented way of implementing file systems.
- VFS allows the same system call interface (the API) to be used for different types of file systems.
- The API is to the VFS interface, rather than any specific type of file system.
- See Figure 11.4.

Virtual File Systems - continued

- Superblock – in memory
 - Device – device identifier
 - Inode pointers
 - **Mounted inode pointer**
 - Covered inode pointer
 - Block size
 - **Superblock operations**
 - A pointer to a set of superblock routines for this file system
 - File system type
 - File system specific

Virtual File Systems - continued

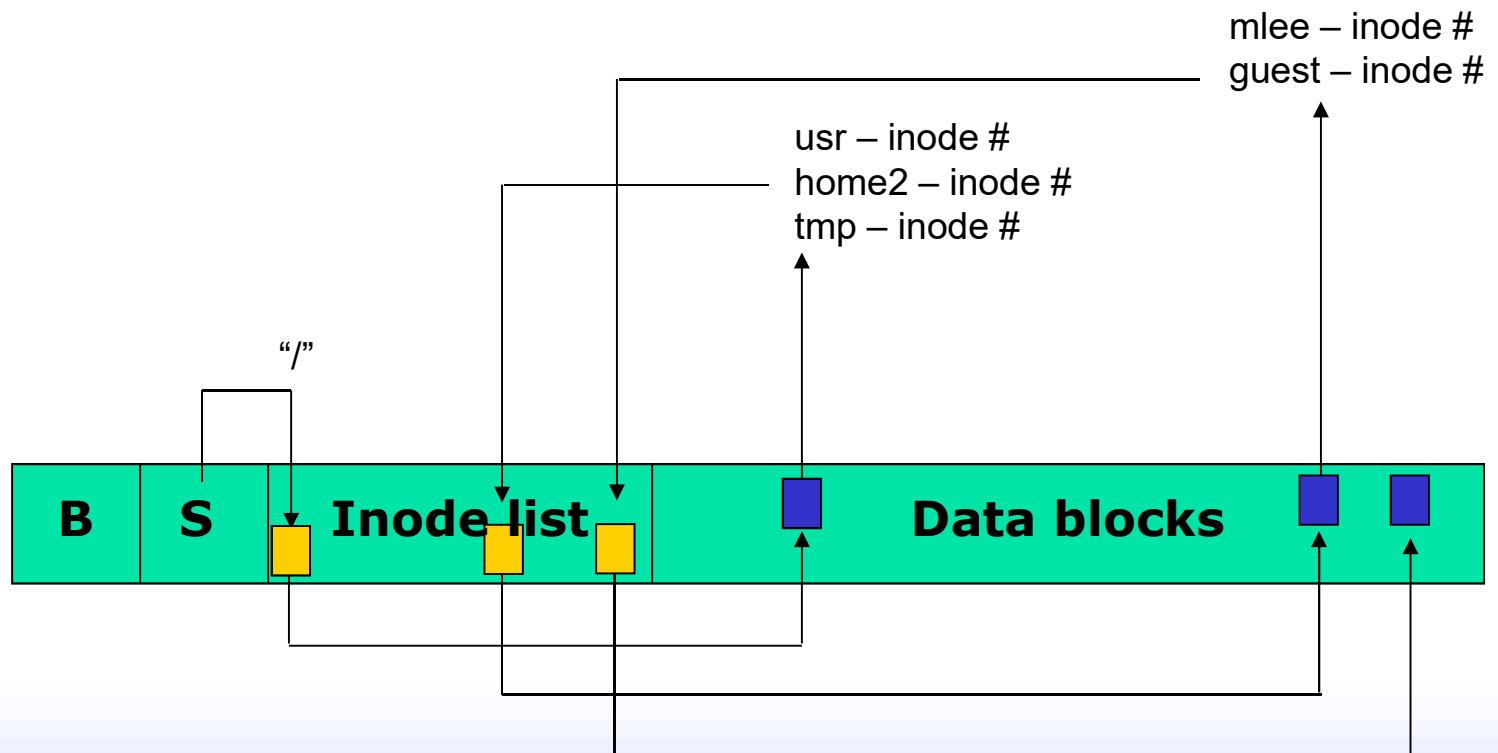
- Inodes for open files and directories – in memory
 - Device – device identifier
 - Inode number
 - The inode number within this file system
 - Mode
 - What it represents
 - User ids
 - The owner identifiers
 - Times
 - Block size
 - Inode operations
 - Count
 - Lock
 - Dirty
 - File system specific information

Virtual File Systems - continued

- Finding a file:
 - From the current working directory – e.g., `$ ls bin`
 - => Inode on the real device for `\.'`
 - File names or directory names with corresponding inodes
 - => Recursive searching ...
 - From the root directory – e.g., `$ ls /usr/src`
 - => Superblock
 - => Device and inode
 - => Inode on the real device for `\/'`
 - File names or directory names with corresponding inodes
 - => Recursive searching...

Virtual File Systems - continued

- Searching a file, e.g., /home2/guest/tst



3. Directory Implementation

- Directory contains a list of file names and corresponding additional information how to locate the file.
- **Linear list** of file names with pointer to the data blocks
 - Advantage
 - Simple to program
 - Disadvantage
 - Time-consuming to execute
- **Hash Table** – linear list with hash data structure
 - Advantage
 - Decreases directory search time
 - Disadvantage
 - *Collisions* – situations where two file names hash to the same location => need to use redundant space
 - Fixed size => hard to grow
- **Sorted list** ...

4. Allocation Methods

- An allocation method refers to how disk blocks are allocated for files:
 - Contiguous allocation
 - Linked allocation
 - Indexed allocation

4.1 Contiguous Allocation

- Each file occupies a set of contiguous blocks on the disk.
- See Figure 11.5.

Block size: 512

File: list

File pointer: 1000

Reading 200 bytes

Block to be accessed:

☺ *Which block[s] would be retrieved from the file system?*

Displacement in the first block to be accessed:

☺ *What displacement is the file pointer in the first block to be accessed?*

Contiguous Allocation - continued

- Advantages
 - Simple – Only starting location (block #) and length (number of blocks) are required.
 - Random access
- Disadvantages
 - Wasteful of space (dynamic storage-allocation problem)
 - => external fragmentation
 - Compaction needed
 - Files cannot grow.

Extent-Based Systems

- Modified contiguous allocation scheme
- Extent-based file systems allocate disk blocks in **extents**.
- An **extent** is a contiguous block of disks. Extents are allocated for file allocation. A file consists of one or more extents.

4.2 Linked Allocation

- Each file is a linked list of disk blocks: blocks may be scattered anywhere on the disk.
- See Figure 11.6.

Block size: 512

Size of the link in a block: 4

File: jeep

File pointer: 1020

Reading 100 bytes

Block to be accessed:

☺ *Which block[s] would be retrieved from the file system?*

Displacement in the first block to be accessed:

☺ *What displacement is the file pointer in the first block to be accessed?*

Linked Allocation - continued

- Advantages
 - Simple – need only starting address
 - Free-space management system – no waste of space
- Disadvantages
 - No random access
 - The space required for the pointers
- A solution
 - Use of clusters of multiple blocks

File-Allocation Table (FAT)

- MS-DOS, OS/2
- See Figure 11.7.
 - FAT: not the list of data blocks
 - The memory of index 339 in FAT is -1.
 - -1 in FAT: End of file

Block size: 512

Size of the link in FAT: 4

File: test

File pointer: 1020

Reading 100 bytes

Block to be accessed:

☺ *Which block[s] would be retrieved from the file system?*

Displacement in the first block to be accessed:

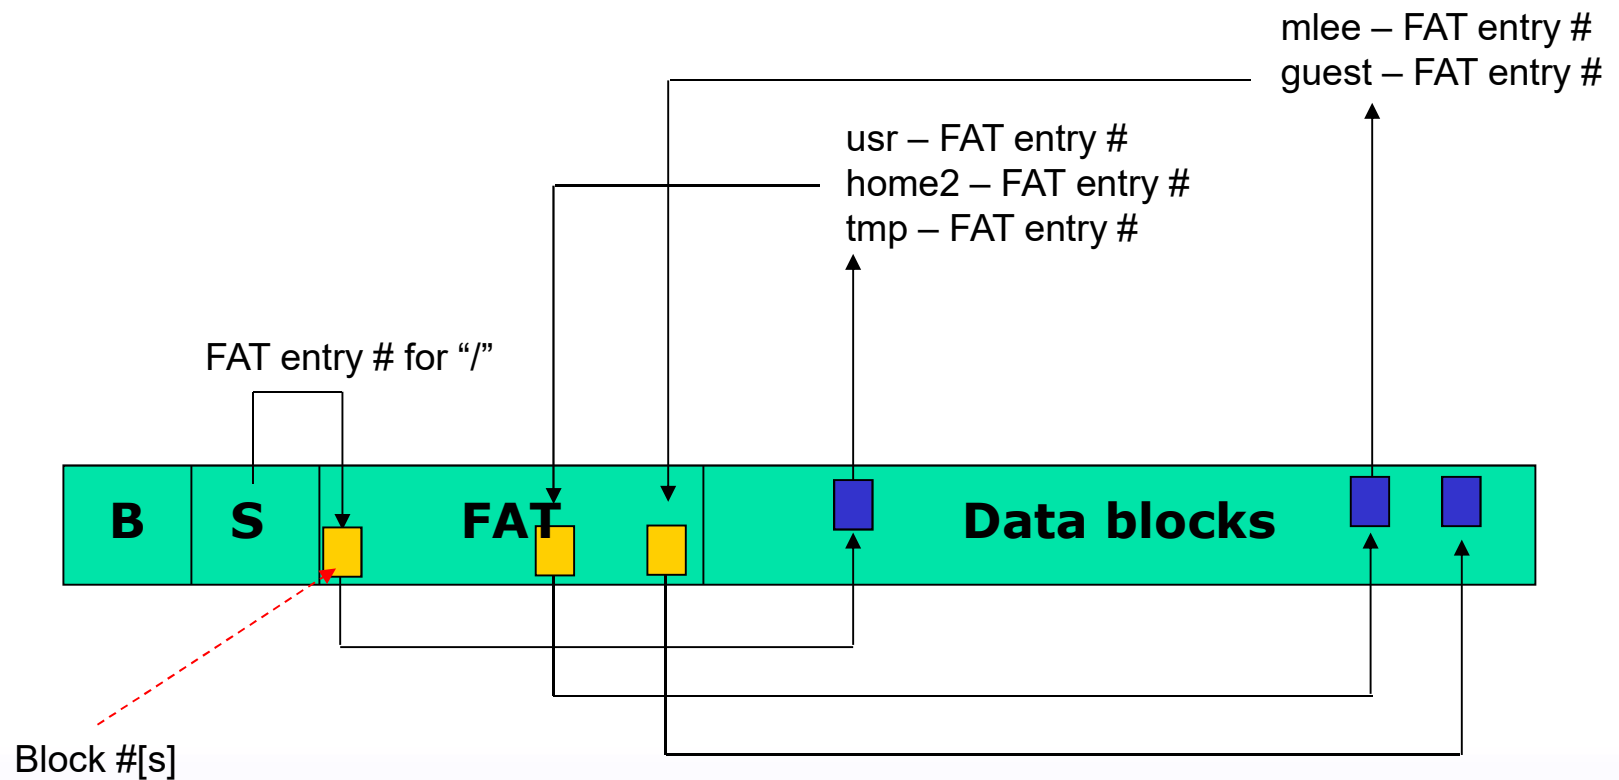
☺ *What displacement is the file pointer in the first block to be accessed?*

☺ *The size of FAT?*

File-Allocation Table - continued

- ☺ *How to allocate a new block to a file?*
 - Another link for unused blocks
- ☺ *Where are directory entry and FAT when the OS is running?*
 - **Directory entry and FAT in main memory to reduce the access time**
- Advantages
 - Actual blocks have no pointer
 - Random access time is improved
- Disadvantages
 - FAT should be in main memory to reduce the access time => difficult to support big file systems

File-Allocation Table - continued



4.3 Indexed Allocation

- Brings all pointers together into the *index block*.
- Logical view
- See Figure 11.8.

Block size: 512

File: jeep

File pointer: 1020

Reading 100 bytes

Block to be accessed:

☺ *Which block[s] would be retrieved from the file system?*

Displacement in the first block to be accessed:

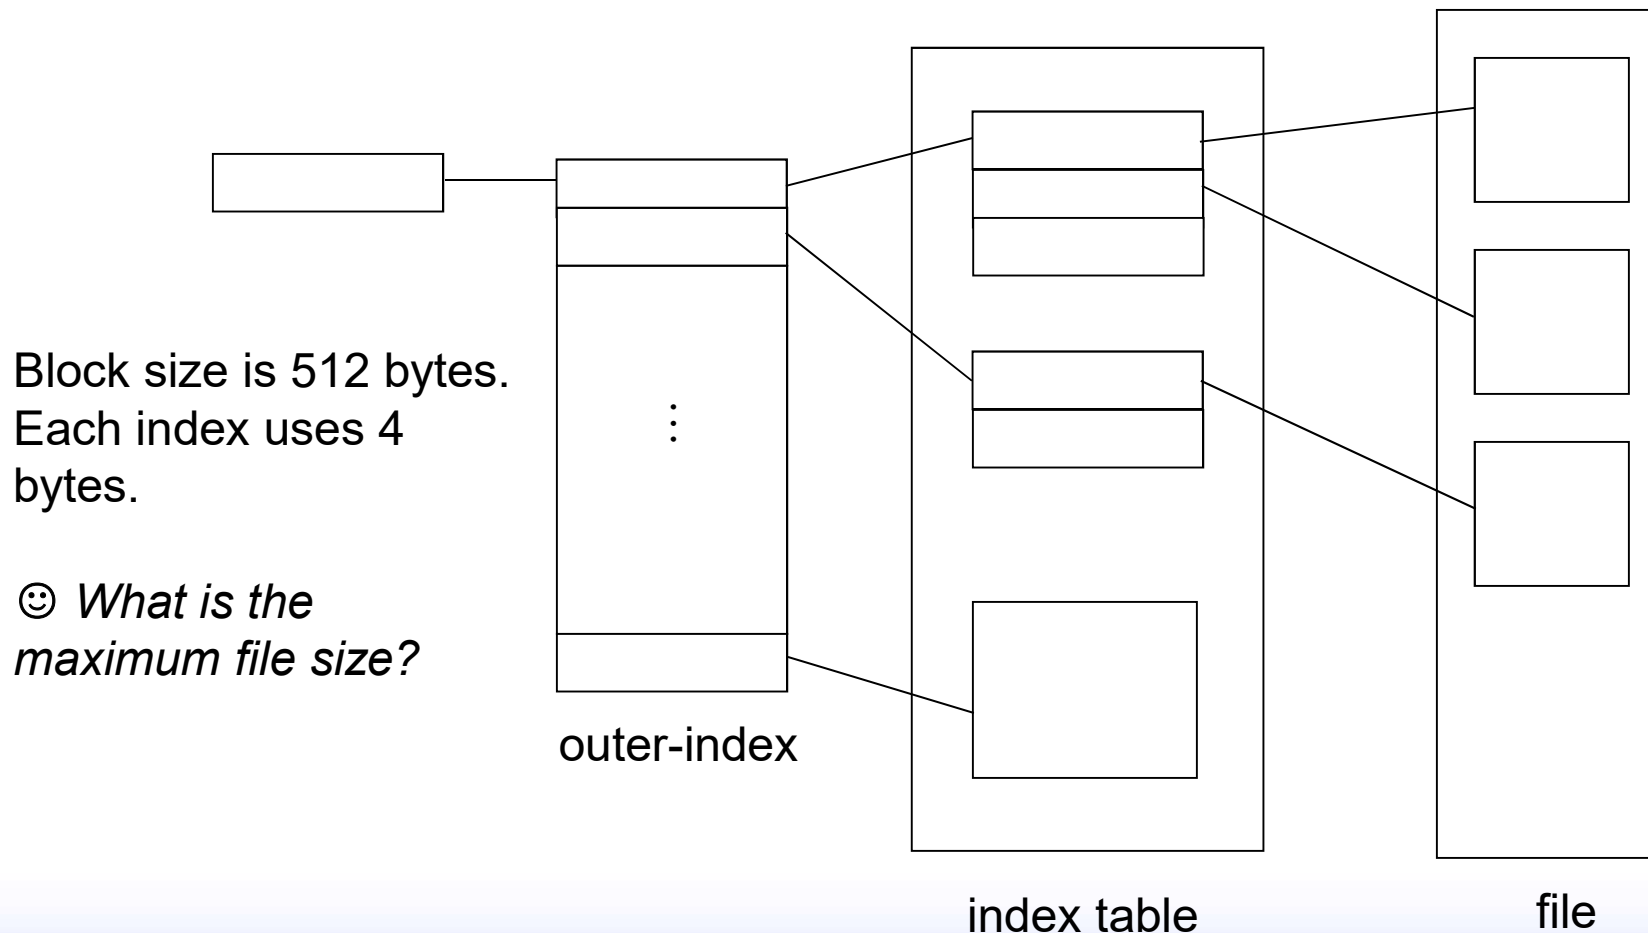
☺ *What displacement is the file pointer in the first block to be accessed?*

☺ *What is the maximum file size when 4 bytes are used for an index?*

Indexed Allocation - continued

- Advantage
 - Random access
 - No external fragmentation
- Disadvantage
 - Need index table

Two-Level Indexed Allocation




Combined Scheme: UNIX (4K bytes per block)

- See Figure 11.9.
 - 10 direct blocks
 - Indirect block: 170 links
 - ☺ *The total addressable space for a file?*

5. Free-Space Management

- **Bit vector** (n blocks)



bit[i] = 0 \Rightarrow block[i] free
 1 \Rightarrow block[i] occupied

- Example

- `int bitVector[512];` // int – 4 bytes
- I would like to see if `block 128` is free.
- ☺ Which element of `bitVector` should be used?

Free-Space Management – cont.

- Advantage
 - Easy to get contiguous files
- Disadvantage
 - Bit map requires extra space [in the main memory]
 - Example:
 - block size = 2^{12} bytes
 - disk size = 2^{30} bytes (1 gigabyte)
 - bit map size $n = 2^{30}/2^{12} = 2^{18}$ bits (or 32K bytes)

Free-Space Management – cont.

- **Linked list** (free list)
 - No waste of space
 - Cannot get contiguous space easily
 - Need to access more block to get several free blocks
 - See Figure 11.10.
- **Linked list in FAT**
- **Grouping**
 - A modification of the free-list approach
 - n free blocks in the first free block
 - The last free block contains the address of another n free blocks, and so on
 - Example: block 2: 3, 4, 5, 8 when four links are allocated to a block
- **Counting**
 - Each entry in the free-space list consists of a block address and a count of contiguous blocks from the block address.
 - Example: block 2: 4, 2; block 4: 8, 2; block 8: 17, 6