

# Chapter 7. Deadlocks


- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock

# The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.
- Example:
  - System has 2 tape drives.
  - $P_1$  and  $P_2$  each hold one tape drive and each needs another one.
- Example:
  - Semaphores  $A$  and  $B$  initialized to 1.

$P_0$	$P_1$
$wait(A);$	$wait(B)$
$wait(B);$	$wait(A)$

# Bridge Crossing Example



- Traffic only in one direction.
- Each section of a bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback).
- Several cars may have to be backed up if a deadlock occurs.
- Starvation is possible.

# 1. System Model

- Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances.
- Each process utilizes a resource as follows:
  - Request
  - Use
  - Release
  - Example of file i/o
    - Open
    - Read/write
    - Close

## 2. Deadlock Characterization

Deadlock can, not will, arise if four conditions hold simultaneously:

- **Mutual exclusion**

- Only one process at a time can use a resource.

- **Hold and wait**

- A process holding at least one resource is waiting to acquire additional resources held by other processes.

- **No preemption**

- A resource can be released only voluntarily by the process holding it, after that process has completed its task.

- **Circular wait**

- There exists a set  $\{P_0, P_1, \dots, P_n, P_0\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .


# Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .

- $V$  is partitioned into two types:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system.
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system.
- Request edge – directed edge  $P_i \rightarrow R_j$
- Assignment edge – directed edge  $R_j \rightarrow P_i$

# Resource-Allocation Graph – cont.


- Process




- Resource type with 4 instances



- $P_i$  requests instance of  $R_j$



- $P_i$  is holding an instance of  $R_j$



# Example of a Resource Allocation Graph

- See Figure 7.1, “Resource-allocation graph”.
  - ☺ *Deadlock?*



# Resource Allocation Graph

- See Figure 7.2, “Resource-allocation graph with a deadlock”.
  - ☺ *Deadlock?*

# Resource Allocation Graph – cont.

- See Figure 7.3, “Resource-allocation graph with a cycle but no deadlock”.
  - ☺ *Deadlock?*

# Basic Facts

- If graph contains no cycles  $\Rightarrow$  no deadlock.
- If graph contains a cycle  $\Rightarrow$ 
  - If only one instance per resource type, then deadlock.
  - If several instances per resource type, possibility of deadlock.

# 3. Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state.
- Allow the system to enter a deadlock state and then recover.
- Ignore the problem and pretend that deadlocks never occur in the system.
  - Used by most operating systems, including UNIX.
  - ☺ *Why?*
- Prevention
- Avoidance
- Detection and recovery

# 4. Deadlock Prevention

Idea: Breaking one of four necessary conditions for deadlocks

- **Mutual Exclusion**

- Not required for sharable resources
- However, must hold for nonsharable resources.

- **Hold and Wait**

- Must guarantee that whenever a process requests a resource, it does not hold any other resources.
- Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none.
- => However, low resource utilization; starvation possible.

# Deadlock Prevention - continued

- **No Preemption**

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
- Preempted resources are added to the list of resources for which the process is waiting.
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.

# Deadlock Prevention - continued

- **Circular Wait**

- $R = \{R_1, R_2, \dots, R_m\}$
- $F: R \rightarrow N$
- **Protocol**
  - Each process requests resources in an increasing order of enumeration.
  - A process can request instances of resource type  $R_j$  iff  $F(R_j) > F(R_i)$  for all  $R_i$  held by the process.
  - If several instances of the same resource type are needed, a single request for all of them must be issued.
  - Then, there will be no circular wait.
- ☺ *Can you prove it?*

- **However, inefficient**

- ☺ *Why???*

# 5. Deadlock Avoidance

Requires that the system has some additional *a priori* information available.

- Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes.



# Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state.
- System is in **safe state** if there exists a **safe sequence** of all processes.
- Sequence  $\langle P_1, P_2, \dots, P_n \rangle$  is **safe** if for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$ .
  - If the resources that  $P_i$  needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished.
  - When all  $P_j$  have finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate.
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on.
  - $\Rightarrow \text{☺ Acyclic graph always?}$

# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks.
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock.
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.
- See Figure 7.8, “Safe, unsafe, and deadlocked state spaces”.

# Resource-Allocation Graph Algorithm

- One instance for each resource type:
  - *Claim edge*  $P_i \rightarrow R_j$  indicated that process  $P_i$  may request resource  $R_j$ ; represented by a dashed line.
  - Claim edge converts to *request edge* when a process requests a resource.
  - When a resource is released by a process, assignment edge reconverts to a claim edge.
  - Resources must be claimed *a priori* in the system.
- See Figure 7.9 – 7.10.

# Banker's Algorithm

- Multiple instances for each resource type:
  - **Each process must a priori claim maximum use.**
  - When a process requests a resource, it may have to wait.
  - When a process gets all its resources, it must return them in a finite amount of time.

# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resource types.

- *Available*: Vector of length  $m$ . If  $Available[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
- *Max*:  $n \times m$  matrix. If  $Max[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
- *Allocation*:  $n \times m$  matrix. If  $Allocation[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .
- *Need*:  $n \times m$  matrix. If  $Need[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.

$$Need[i,j] = Max[i,j] - Allocation[i,j].$$

# Safety Algorithm

Find out whether or not a system is in a safe state:

**Try to find a safe sequence of all processes.**

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:

*Work* = *Available*

- resources

*Finish*[*i*] = *false* for *i* = 0, 1, 2, 3, ..., *n*-1.

- processes

2. Find an *i* such that both:

(a) *Finish*[*i*] = *false*

(b)  $Need_i \leq Work$

If no such *i* exists, go to step 4.

3.  $Work = Work + Allocation_i$

*Finish*[*i*] = *true*

go to step 2.

4. If *Finish*[*i*] == *true* for all *i*, then the system is in a safe state.

# Resource-Request Algorithm for Process $P_i$

$Request_i$  = request vector for process  $P_i$ . If  $Request_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$ .

1. If  $Request_i \leq Need_i$  go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If  $Request_i \leq Available$ , go to step 3. Otherwise  $P_i$  must wait, since resources are not available.
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If safe, i.e., there is a safe sequence  $\Rightarrow$  the resources are allocated to  $P_i$ .
- If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

**However, inefficient**

☺ **Why???**

**The above operation is very expensive because of high complexity,  $O(mn^2)$ , of the safety checking algorithm.**

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$ ; 3 resource types
  - $A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances).
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	



# Example - continued

- The content of the matrix. Need is defined to be Max – Allocation.

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>	<u>Need</u>
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
$P_0$	0 1 0	7 5 3	3 3 2	7 4 3
$P_1$	2 0 0	3 2 2		1 2 2
$P_2$	3 0 2	9 0 2		6 0 0
$P_3$	2 1 1	2 2 2		0 1 1
$P_4$	0 0 2	4 3 3		4 3 1

- ☺ *The system is in a safe state?*
- The sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria.

# Example $P_1$ Request (1,0,2) Cont.

- Check that  $\text{Request}_1 (1,0,2) \leq \text{Available}$  (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$ )

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement.

# Example $P_1$ Request (1,0,2) Cont.

Currently

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- ☺ Can request for (3,3,0) by  $P_4$  be granted?
  - No, because  $Need_4 > Available_4$
- ☺ Can request for (0,2,0) by  $P_0$  be granted?
  - No, because  $Available:(2,3,0) \Rightarrow (2,1,0)$ , and then no process can finish.

# 6. Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

# Single Instance of Each Resource Type

- Maintain *wait-for* graph
  - Nodes are processes.
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$ .
- Periodically invoke an **algorithm that searches for a cycle** in the graph.
- An algorithm to detect a cycle in a graph requires  $O(n^2)$  operations, where  $n$  is the number of vertices in the graph.
- See Figure 7.11.

# Several Instances of a Resource Type

- *Available:* A vector of length  $m$  indicates the number of available resources of each type.
- *Allocation:* An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
- *Request:* An  $n \times m$  matrix indicates the current request of each process. If  $Request_i[j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Detection Algorithm

Try to find a safe sequence of all processes

1. Let  $Work$  and  $Finish$  be vectors of length  $m$  and  $n$ , respectively Initialize:

(a)  $Work = Available$

(b) For  $i = 1, 2, \dots, n$ ,

if  $Allocation_i \neq 0$ , then

$Finish[i] = false$ ;

otherwise,

$Finish[i] = true$ .

2. Find an index  $i$  such that both:

(a)  $Finish[i] == false$

(b)  $Request_i \leq Work$

If no such  $i$  exists, go to step 4.

# Detection Algorithm - continued

3.  $Work = Work + Allocation_i$   
 $Finish[i] = true$   
go to step 2.
4. If  $Finish[i] == false$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state.  
Moreover, if  $Finish[i] == false$ , then  $P_i$  is deadlocked.

Algorithm requires an order of  $O(mn^2)$  operations to detect whether the system is in deadlocked state.



# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types  $A$  (7 instances),  $B$  (2 instances), and  $C$  (6 instances).
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
$P_0$	0 1 0	0 0 0	0 0 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 0	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

- Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$ .

# Example - continued

- $P_2$  requests an additional instance of type C.

	<u>Request</u>		
	A	B	C
$P_0$	0	0	0
$P_1$	2	0	2
$P_2$	0	0	1
$P_3$	1	0	0
$P_4$	0	0	2

- ☺ *State of system?*
  - Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes; requests.
  - Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$ .

# Detection-Algorithm Usage

- ☺ *When, and how often, to invoke???*
- It depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - One for each disjoint cycle
- If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph,
- And we would not be able to tell which of the many deadlocked processes “caused” the deadlock.

## 7. Recovery from Deadlock: Process Termination

- Abort all deadlocked processes.
- Abort one process at a time until the deadlock cycle is eliminated.
- ☺ *In which order should we choose to abort?*
  - Priority of the process.
  - How long the process has computed, and how much longer to completion.
  - Resources the process has used.
  - Resources the process needs to complete.
  - How many processes will need to be terminated?
  - ☺ *Is process interactive or batch?*

# Recovery from Deadlock: Resource Preemption

- Selecting a victim – minimize cost.
- Rollback – return to some safe state, restart process for that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.

# Combined Approach to Deadlock Handling

- Combine the three basic approaches:

- Prevention
- Avoidance
- Detection

allowing the use of the optimal approach for each of resources in the system.

- Partition resources into hierarchically ordered classes.
- Use most appropriate technique for handling deadlocks within each class.